

Detecting Driver Drowsiness Using Image Processing Techniques

Nikhil Mishra - 210668, Dhruv - 210338, Rishi Poonia - 210851
IIT Kanpur

November 9, 2024

Abstract

Driver drowsiness is a critical factor contributing to road accidents worldwide. This project presents an innovative approach to detecting driver drowsiness through image processing by analyzing facial landmarks. Utilizing eye aspect ratio, pupil-to-eye center distance, mouth aspect ratio, and mouth-to-eye ratio, we extract vital features from video frames to identify signs of drowsiness. These features are then processed using a Long Short-Term Memory (LSTM) network to classify the driver's state. Our method demonstrates high accuracy in distinguishing between drowsy and alert drivers, offering a promising solution for enhancing road safety.

1 Introduction

Driver drowsiness is a major contributor to vehicular accidents worldwide, resulting in numerous severe injuries and fatalities each year. Traditional methods for detecting drowsiness often rely on physiological sensors or vehicle-based systems. While effective to some extent, these approaches can be intrusive, uncomfortable for drivers, and limited in their scope of functionality.

In contrast, image processing offers a non-intrusive and scalable solution by analyzing facial cues that indicate fatigue. This project leverages advanced image processing techniques to monitor driver alertness in real-time. By extracting and examining facial landmarks, with a specific focus on the eyes and mouth, we aim to identify patterns associated with drowsiness. The novelty of our approach lies in the combination of multiple facial ratios and the application of deep learning models, which together enhance the accuracy and reliability of drowsiness detection.

2 Proposed Method

Our methodology encompasses several key steps: video preprocessing, feature extraction, data segmentation, and classification using an LSTM network. The overall framework is illustrated in Figure 1.

2.1 Video Preprocessing

Each video is divided into frames, and facial landmarks are detected using the **dlib** library. Key facial regions, including the eyes and mouth, are identified to extract relevant features.

2.2 Feature Extraction

For each frame extracted from the video, we calculate four primary ratios that are indicative of driver drowsiness: Eye Aspect Ratio (EAR), Pupil-to-Eye Center Distance (PUC), Mouth Aspect Ratio (MAR), and Mouth-to-Eye Ratio (MOE). These ratios are derived from the coordinates of specific facial landmarks (Figure 2.) obtained through facial landmark detection.

2.2.1 Eye Aspect Ratio (EAR)

Definition: The Eye Aspect Ratio (EAR) is a measure of the openness of the eyes. It is designed to remain approximately constant when the eyes are open and rapidly decrease to zero during a blink.

Formula:

$$\text{EAR} = \frac{||p_2 - p_6|| + ||p_3 - p_5||}{2 \times ||p_1 - p_4||} \quad (1)$$

Where:

- $p_1, p_2, p_3, p_4, p_5, p_6$ are the 2D coordinates of the six key eye landmarks, typically corresponding to the outer corner, upper eyelid, lower eyelid, and inner corner of the eye.
- $||p_i - p_j||$ denotes the Euclidean distance between points p_i and p_j .

Rationale: The EAR captures the ratio of vertical eye landmarks to the horizontal eye landmarks. When the eye is open, the vertical distances $||p_2 - p_6||$ and $||p_3 - p_5||$ maintain a certain proportion relative to the horizontal distance $||p_1 - p_4||$. During a blink or prolonged eye closure, the vertical distances decrease sharply while the horizontal distance remains relatively unchanged, causing the EAR to drop rapidly. This characteristic makes EAR a reliable indicator for detecting blinks and sustained eye closures, both of which are associated with drowsiness.

2.2.2 Mouth Aspect Ratio (MAR)

Definition: The Mouth Aspect Ratio (MAR) quantifies the openness of the mouth, which can be indicative of yawning—a common sign of drowsiness.

Formula:

$$\text{MAR} = \frac{||p_{13} - p_{19}|| + ||p_{14} - p_{18}|| + ||p_{15} - p_{17}||}{3 \times ||p_{12} - p_{16}||} \quad (2)$$

Rationale: Similar to EAR, MAR captures the ratio of vertical mouth landmarks to the horizontal mouth landmarks. When the mouth is open, such as during yawning, the vertical distance increases significantly relative to the horizontal distance

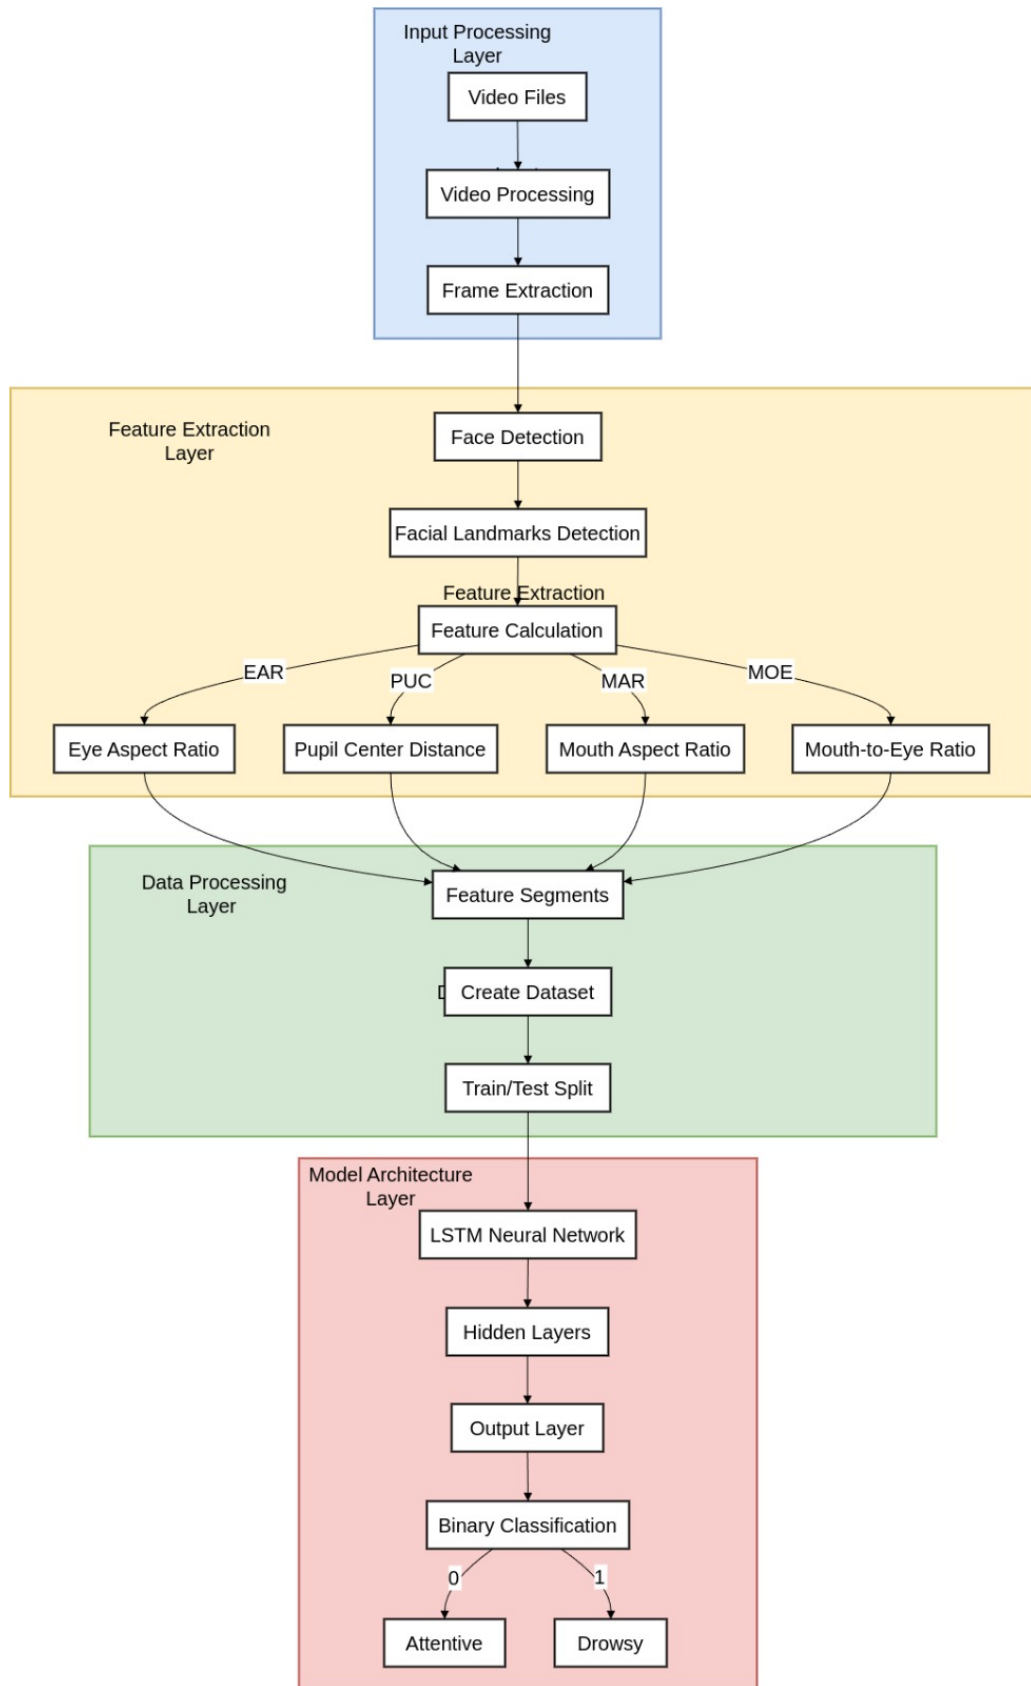


Figure 1: Overview of the Proposed Drowsiness Detection Method



Figure 2: The indexes of the 68-coordinates corresponding to the facial landmarks




Figure 3: Left: A visualization of eye landmarks when the eye is open. Right: Eye landmarks when the eye is closed

2.2.3 Pupil-to-Eye Center Distance (PUC)

Definition: The Pupil-to-Eye Center Distance (PUC) measures the distance between the center of the pupil and the geometric center of the eye.

Formula:

$$\text{PUC} = ||p_{\text{pupil}} - p_{\text{eye center}}|| \quad (3)$$

Where:

- p_{pupil} is the 2D coordinate of the pupil.
- $p_{\text{eye center}}$ is the 2D coordinate of the geometric center of the eye, typically calculated as the midpoint between the inner and outer corners of the eye.

2.2.4 Mouth-to-Eye Ratio (MOE)

Definition: The Mouth-to-Eye Ratio (MOE) is the ratio of the Mouth Aspect Ratio (MAR) to the Eye Aspect Ratio (EAR). It provides a combined measure of the mouth and eye states, encapsulating both yawning and eye closure behaviors.

Formula:

$$\text{MOE} = \frac{\text{MAR}}{\text{EAR}} \quad (4)$$

2.3 Data Segmentation

Videos are segmented into fixed-length frames (e.g., 150 frames per segment). For each segment, a feature vector is created by compiling the four ratios across all frames. Each segment is labeled as drowsy (1) or alert (0) based on the video source.

2.4 Classification with LSTM

The segmented feature vectors are fed into a Long Short-Term Memory (LSTM) network, which is adept at handling sequential data. The LSTM model learns temporal dependencies in the facial feature sequences to accurately classify the driver’s state.


3 Experimental Results

To evaluate the effectiveness of our approach, we conducted experiments using the UTA RealLife Drowsiness Dataset. The dataset comprises videos labeled as drowsy or alert, providing a robust basis for training and testing our model.


The model achieved the following performance metrics:

	precision	recall	f1-score	support
0	0.82	0.82	0.82	114
1	0.78	0.78	0.78	94
accuracy	0.80			208
macro avg	0.80	0.80	0.80	208
weighted avg	0.80	0.80	0.80	208

Table 1: Classification Report of Best Model



(a) Confusion Matrix for Best Model



(b) ROC Curve for Best Model

Figure 4: Performance Metrics of the Best Model

4 Experimental Analysis

In this section, we analyze the performance of two neural network architectures—GRUNet and LightCNNLSTM—used for driver drowsiness detection.

4.1 GRUNet

GRUNet is a bidirectional Gated Recurrent Unit (GRU)-based network designed for sequence classification. It consists of two bidirectional GRU layers with dropout for regularization and batch normalization to stabilize training. The final hidden states are passed through fully connected layers with ReLU activation and dropout, culminating in a two-class softmax output. GRUNet effectively captures temporal dependencies in the facial feature sequences, enabling reliable classification of drowsy and alert states.

4.2 LightCNNLSTM

LightCNNLSTM integrates one-dimensional Convolutional Neural Networks (1D CNNs) with bidirectional Long Short-Term Memory (LSTM) layers to harness both spatial and temporal feature extraction. The architecture begins with convolutional and max-pooling layers to extract local temporal features, followed by bidirectional LSTM layers that model long-term dependencies. Fully connected layers with batch normalization, ReLU activation, and dropout lead to a two-class softmax output. This hybrid approach enhances feature representation, improving classification accuracy for drowsiness detection.

4.3 Comparison of Models

Both models are evaluated on the same dataset to ensure a fair comparison. GRUNet offers computational efficiency with its GRU layers, while LightCNNLSTM leverages the combined strengths of CNNs and LSTMs for enhanced feature extraction and sequence modeling. The inclusion of convolutional layers in LightCNNLSTM allows for better spatial feature capture, resulting in superior performance metrics.




Figure 5: Confusion Matrix For GRUNet and LightCNNLstm

5 Conclusion

This project demonstrates an effective method for detecting driver drowsiness through image processing and deep learning. By analyzing facial landmarks and leveraging temporal patterns with an LSTM network, our approach achieves high accuracy in distinguishing between drowsy and alert drivers. Future work may explore real-time implementation and integration with vehicle systems to enhance road safety further.

References

- [1] Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735-1780.
- [2] University of Texas at Arlington. UTA RealLife Drowsiness Dataset. <https://www.uta.edu/datasets/drowsiness>

A Appendix

A.1 Preprocessing Code

```
1 import os
2 import cv2
3 import dlib
4 import numpy as np
5 import pickle
6
7 def eye_aspect_ratio(eye):
8     x = [point.x for point in eye]
9     y = [point.y for point in eye]
10    A = np.linalg.norm(np.array([x[1] - x[5], y[1] - y[5]]))
11    B = np.linalg.norm(np.array([x[2] - x[4], y[2] - y[4]]))
12    C = np.linalg.norm(np.array([x[0] - x[3], y[0] - y[3]]))
13    ear = (A + B) / (2.0 * C)
14    return ear
15
16 def pupil_to_eye_center_distance(eye):
17     x = [point.x for point in eye]
18     y = [point.y for point in eye]
19     d = np.linalg.norm(np.array([x[0] - x[3], y[0] - y[3]]))
20     return d
21
22 def mouth_aspect_ratio(mouth):
23     x = [point.x for point in mouth]
24     y = [point.y for point in mouth]
25     A = np.linalg.norm(np.array([x[13] - x[19], y[13] - y[19]]))
26     B = np.linalg.norm(np.array([x[14] - x[18], y[14] - y[18]]))
27     C = np.linalg.norm(np.array([x[15] - x[17], y[15] - y[17]]))
28     mar = (A + B + C) / (3.0 * np.linalg.norm(np.array([x[12] - x[16], y[12]
29     - y[16]])))
30     return mar
31
32 def mouth_to_eye_ratio(eye, mouth):
33     ear = eye_aspect_ratio(eye)
34     mar = mouth_aspect_ratio(mouth)
35     if ear == 0:
36         return 0
37     moe = mar / ear
38     return moe
39
40 def extract_features(frame, detector, predictor):
41     gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
42     faces = detector(gray)
43     features = []
44
45     for face in faces:
46         shape = predictor(gray, face)
47         ear = eye_aspect_ratio(shape.parts()[36:42])
48         puc = pupil_to_eye_center_distance(shape.parts()[36:42])
49         mar = mouth_aspect_ratio(shape.parts()[48:68])
50         moe = mouth_to_eye_ratio(shape.parts()[36:42], shape.parts()[48:68])
51         features.append([ear, puc, mar, moe])
52
53     return features[0] if features else [0, 0, 0, 0]
54
55 def process_video_segments(video_path, detector, predictor,
56                             frames_per_segment=50, num_segments=2):
```



```

55     cap = cv2.VideoCapture(video_path)
56     total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
57
58     if total_frames < frames_per_segment:
59         cap.release()
60         return []
61
62     segment_positions = np.linspace(0, total_frames - frames_per_segment,
63 num_segments, dtype=int)
64     features_segments = []
65
66     ctr = 0
67
68     for start_pos in segment_positions:
69
70         ctr += 1
71         print(ctr)
72
73         cap.set(cv2.CAP_PROP_POS_FRAMES, start_pos)
74         current_segment = []
75
76         for _ in range(frames_per_segment):
77             ret, frame = cap.read()
78             if not ret:
79                 break
80             features = extract_features(frame, detector, predictor)
81             current_segment.append(features)
82
83             if len(current_segment) == frames_per_segment:
84                 features_segments.append(current_segment)
85                 print(current_segment)
86
87     cap.release()
88     return features_segments
89
90 # Initialize face detection tools
91 detector = dlib.get_frontal_face_detector()
92 predictor = dlib.shape_predictor("/kaggle/input/shape-predictor-68-face-
93 landmarksdat/shape_predictor_68_face_landmarks.dat")
94
95 base_path = "/kaggle/input/uta-reallife-drowsiness-dataset/"
96 frames_per_segment = 150
97 num_segments_per_video = 20
98 output_file = "drowsiness_features.pkl"
99
100 all_features = []
101 all_labels = []
102
103 # Repeat multiple times, changing folder paths to Fold1_part2/Fold1_part2,
104 etc.
105 folder_path = os.path.join(base_path, "Fold1_part1/Fold1_part1")
106 for subject_folder in os.listdir(folder_path):
107     subject_path = os.path.join(folder_path, subject_folder)
108     print(f"Processing subject: {subject_folder}")
109
110     for video_file in os.listdir(subject_path):
111         if video_file in ['0.mov', '10.mov', '0.MOV', '10.MOV', '0.mp4', '
112 10.mp4']:

```

```

111         video_path = os.path.join(subject_path, video_file)
112         label = 1 if video_file.startswith('10') else 0
113         print(f"Processing video: {video_file}")
114
115         segments = process_video_segments(video_path, detector,
116 predictor,
117                                     frames_per_segment,
118 num_segments_per_video)
119
120         all_features.extend(segments)
121         all_labels.extend([label] * len(segments))
122
123 print(f"Total segments collected: {len(all_features)}")
124
125 # Save features and labels
126 data = {
127     'features': np.array(all_features),
128     'labels': np.array(all_labels)
129 }
130
131 with open(output_file, 'wb') as f:
132     pickle.dump(data, f)
133
134 print(f"Features and labels saved to {output_file}")

```

Listing 1: Preprocessing Code

A.2 Model Training Code

```

1 import torch
2 import torch.nn as nn
3 from torch.utils.data import Dataset, DataLoader
4 import pickle
5 import numpy as np
6 from sklearn.model_selection import train_test_split
7 from sklearn.metrics import confusion_matrix, classification_report,
8   roc_curve, auc
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from copy import deepcopy
12
13 class DrowsinessDataset(Dataset):
14     def __init__(self, features, labels):
15         self.features = torch.FloatTensor(features)
16         self.labels = torch.LongTensor(labels)
17
18     def __len__(self):
19         return len(self.labels)
20
21     def __getitem__(self, idx):
22         return self.features[idx], self.labels[idx]
23
24 class DrowsinessLSTM(nn.Module):
25     def __init__(self, input_size=4, hidden_size=128, num_layers=2,
26 bidirectional=True, dropout=0.5):
27         super(DrowsinessLSTM, self).__init__()
28         self.hidden_size = hidden_size
29         self.num_layers = num_layers
30         self.bidirectional = bidirectional

```

```

29     self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first
= True,
30                           bidirectional=bidirectional, dropout=dropout)
31
32     fc_input_size = hidden_size * (2 if bidirectional else 1)
33
34     self.fc1 = nn.Linear(fc_input_size, fc_input_size // 2)
35     self.fc2 = nn.Linear(fc_input_size // 2, 2)
36
37     self.dropout = nn.Dropout(dropout)
38     self.batch_norm = nn.BatchNorm1d(fc_input_size // 2)
39
40     def forward(self, x):
41         h0 = torch.zeros(self.num_layers * (2 if self.bidirectional else 1),
42                           x.size(0), self.hidden_size).to(x.device)
43         c0 = torch.zeros(self.num_layers * (2 if self.bidirectional else 1),
44                           x.size(0), self.hidden_size).to(x.device)
45
46         out, _ = self.lstm(x, (h0, c0))
47         out = self.fc1(out[:, -1, :])
48         out = self.batch_norm(out)
49         out = torch.relu(out)
50         out = self.dropout(out)
51         out = self.fc2(out)
52
53         return out
54
55     def plot_metrics(history):
56         """Plot training history metrics."""
57         fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))
58
59         # Plot training and validation loss
60         ax1.plot(history['train_loss'], label='Training Loss')
61         ax1.plot(history['val_loss'], label='Validation Loss')
62         ax1.set_title('Model Loss Over Time')
63         ax1.set_xlabel('Epoch')
64         ax1.set_ylabel('Loss')
65         ax1.legend()
66
67         # Plot validation accuracy
68         ax2.plot(history['val_acc'], label='Validation Accuracy')
69         ax2.set_title('Validation Accuracy Over Time')
70         ax2.set_xlabel('Epoch')
71         ax2.set_ylabel('Accuracy')
72         ax2.legend()
73
74         plt.tight_layout()
75         plt.show()
76
77     def plot_confusion_matrix(true_labels, predictions):
78         cm = confusion_matrix(true_labels, predictions)
79         plt.figure(figsize=(8, 6))
80         sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
81         plt.title('Confusion Matrix')
82         plt.ylabel('True Label')
83         plt.xlabel('Predicted Label')
84         plt.show()
85
86     def plot_roc_curve(true_labels, pred_probs):
87         fpr, tpr, _ = roc_curve(true_labels, pred_probs[:, 1])

```

```

88     roc_auc = auc(fpr, tpr)
89
90     plt.figure(figsize=(8, 6))
91     plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
92     plt.plot([0, 1], [0, 1], 'k--')
93     plt.xlim([0.0, 1.0])
94     plt.ylim([0.0, 1.05])
95     plt.xlabel('False Positive Rate')
96     plt.ylabel('True Positive Rate')
97     plt.title('Receiver Operating Characteristic (ROC) Curve')
98     plt.legend(loc="lower right")
99     plt.show()
100
101 def evaluate_model(model, test_loader, device):
102     model.eval()
103     all_preds = []
104     all_labels = []
105     all_pred_probs = []
106
107     with torch.no_grad():
108         for features, labels in test_loader:
109             features, labels = features.to(device), labels.to(device)
110             outputs = model(features)
111             pred_probs = torch.softmax(outputs, dim=1)
112             _, predicted = torch.max(outputs.data, 1)
113
114             all_preds.extend(predicted.cpu().numpy())
115             all_labels.extend(labels.cpu().numpy())
116             all_pred_probs.extend(pred_probs.cpu().numpy())
117
118     all_preds = np.array(all_preds)
119     all_labels = np.array(all_labels)
120     all_pred_probs = np.array(all_pred_probs)
121
122     print("\nClassification Report:")
123     print(classification_report(all_labels, all_preds))
124
125     plot_confusion_matrix(all_labels, all_preds)
126
127     plot_roc_curve(all_labels, all_pred_probs)
128
129     return all_preds, all_labels, all_pred_probs
130
131 def train_model(model, train_loader, test_loader, criterion, optimizer,
132                 num_epochs, device):
133     best_val_acc = 0
134     best_model = None
135     history = {
136         'train_loss': [],
137         'val_loss': [],
138         'val_acc': []
139     }
140
141     for epoch in range(num_epochs):
142         model.train()
143         total_loss = 0
144         for batch_features, batch_labels in train_loader:
145             batch_features, batch_labels = batch_features.to(device),

```

```

146         optimizer.zero_grad()
147         outputs = model(batch_features)
148         loss = criterion(outputs, batch_labels)
149         loss.backward()
150         optimizer.step()
151
152         total_loss += loss.item()
153
154     avg_loss = total_loss / len(train_loader)
155     history['train_loss'].append(avg_loss)
156
157     # Validation
158     model.eval()
159     correct = 0
160     total = 0
161     val_loss = 0
162
163     with torch.no_grad():
164         for features, labels in test_loader:
165             features, labels = features.to(device), labels.to(device)
166             outputs = model(features)
167             loss = criterion(outputs, labels)
168             val_loss += loss.item()
169             _, predicted = torch.max(outputs.data, 1)
170             total += labels.size(0)
171             correct += (predicted == labels).sum().item()
172
173     val_accuracy = 100 * correct / total
174     avg_val_loss = val_loss / len(test_loader)
175     history['val_loss'].append(avg_val_loss)
176     history['val_acc'].append(val_accuracy)
177
178     print(f'Epoch [{epoch+1}/{num_epochs}]')
179     print(f'Training Loss: {avg_loss:.4f}')
180     print(f'Validation Loss: {avg_val_loss:.4f}')
181     print(f'Validation Accuracy: {val_accuracy:.2f}%')
182
183     # Save best model
184     if val_accuracy > best_val_acc:
185         best_val_acc = val_accuracy
186         best_model = deepcopy(model.state_dict())
187         print(f'New best model saved with validation accuracy: {
188 val_accuracy:.2f}%')
189         print('-' * 60)
190
191     # Plot training history
192     plot_metrics(history)
193
194     return best_model, history
195
196 def main():
197     # Load the preprocessed data
198     with open('/kaggle/input/combined-features/combined_features.pkl', 'rb')
199     as f:
200         data = pickle.load(f)
201
202     X = data['features']
203     y = data['labels']
204
205     # Split data

```

```

204 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
205 random_state=42)
206
207 # Parameters
208 batch_size = 500
209 num_epochs = 1000
210 learning_rate = 0.001
211
212 # Create datasets and dataloaders
213 train_dataset = DrowsinessDataset(X_train, y_train)
214 test_dataset = DrowsinessDataset(X_test, y_test)
215
216 train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=
True)
217 test_loader = DataLoader(test_dataset, batch_size=batch_size)
218
219 # Initialize model, loss function, and optimizer
220 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
221 model = DrowsinessLSTM(input_size=4, hidden_size=128, num_layers=2,
222                        bidirectional=True, dropout=0.5).to(device)
223 criterion = nn.CrossEntropyLoss()
224 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
225
226 # Train the model and get the best version
227 best_model_state, history = train_model(model, train_loader, test_loader
,
228                                         criterion, optimizer, num_epochs,
229                                         device)
230
231 # Load the best model
232 model.load_state_dict(best_model_state)
233
234 print("\nEvaluating Model:")
235 predictions, true_labels, pred_probs = evaluate_model(model, test_loader
, device)
236
237 # Save the best model
238 torch.save({
239     'model_state_dict': best_model_state,
240     'history': history,
241     'test_predictions': predictions,
242     'test_labels': true_labels,
243     'test_probabilities': pred_probs
244 }, 'best_drowsiness_model.pth')
245 print("\nModel and metrics saved successfully!")
246
247 if __name__ == "__main__":
248     main()

```

Listing 2: Model Training Code

A.3 Dataset Links

- UTA RealLife Drowsiness Dataset: <https://paperswithcode.com/dataset/uta-rldd>